



INDRAPRASTHA INSTITUTE *of*
INFORMATION TECHNOLOGY **DELHI**

ACM India Winter School on Full-stack Networking

Day 1

Understanding Linux Networking Concepts

Rinku Shah, Assistant Professor, IIIT Delhi

Sumit Kumar, Research Scholar, IIIT Delhi

Najiya Naj, Research Scholar, IIIT Delhi

Maryam Tahira, Research Scholar, IIIT Delhi

Brief Bio



Rinku Shah
Assistant Professor, IIITD
rinku@iiitd.ac.in
<https://faculty.iiitd.ac.in/~rinku/>



Sumit Kumar
Research Scholar, IIITD
sumitk@iiitd.ac.in



Najiya Naj
Research Scholar, IIITD
najiyan@iiitd.ac.in

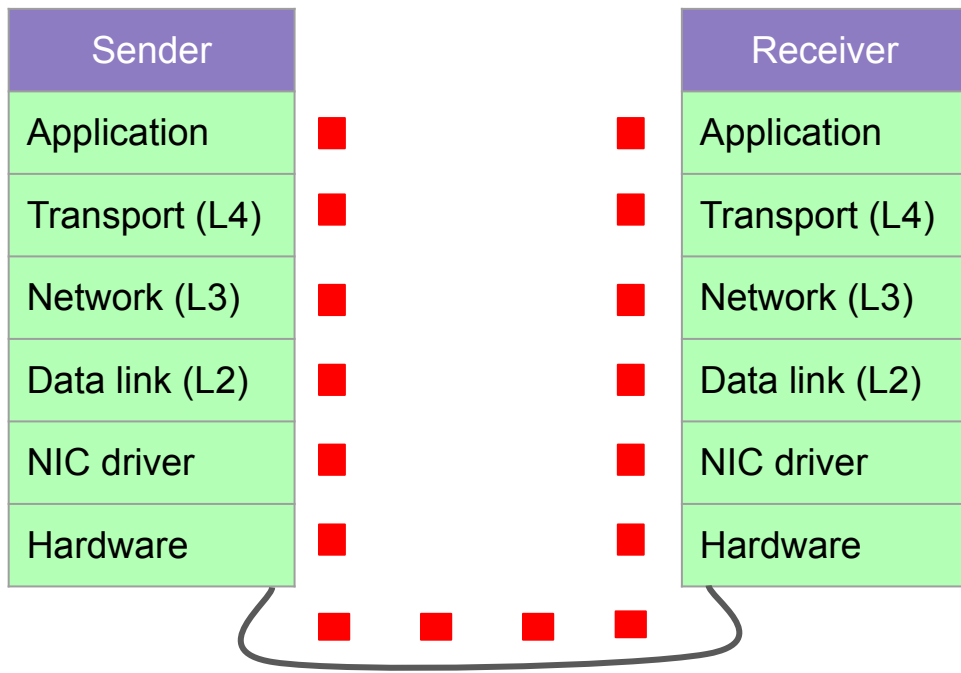


Maryam Tahira
Research Scholar, IIITD
maryamt@iiitd.ac.in

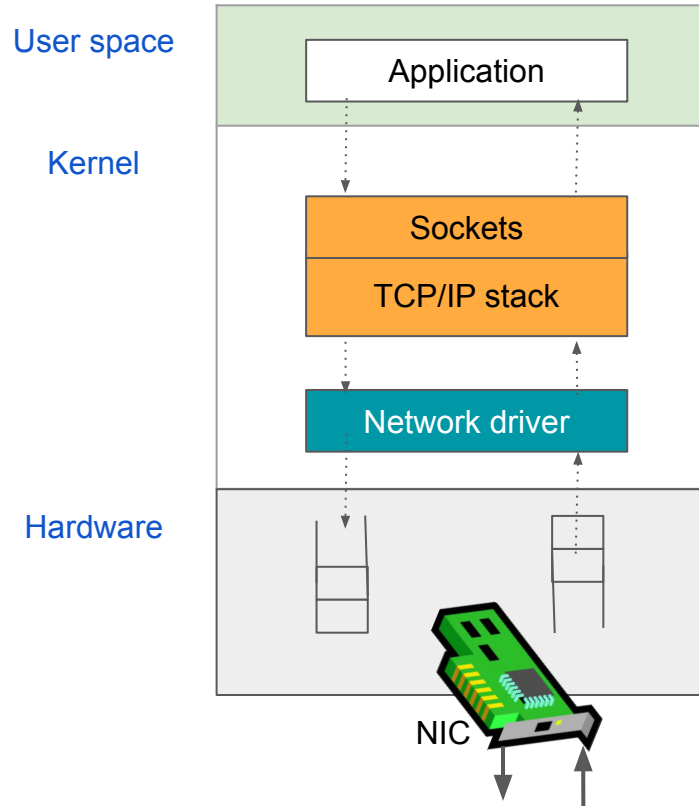
Outline

- Linux networking stack
 - Journey of a packet
- Evolution of NICs
- Evolution of network packet processors
- Hands-on session with Linux system
 - Basic networking commands
 - Configuring network functions

TCP packet flow from source to destination application

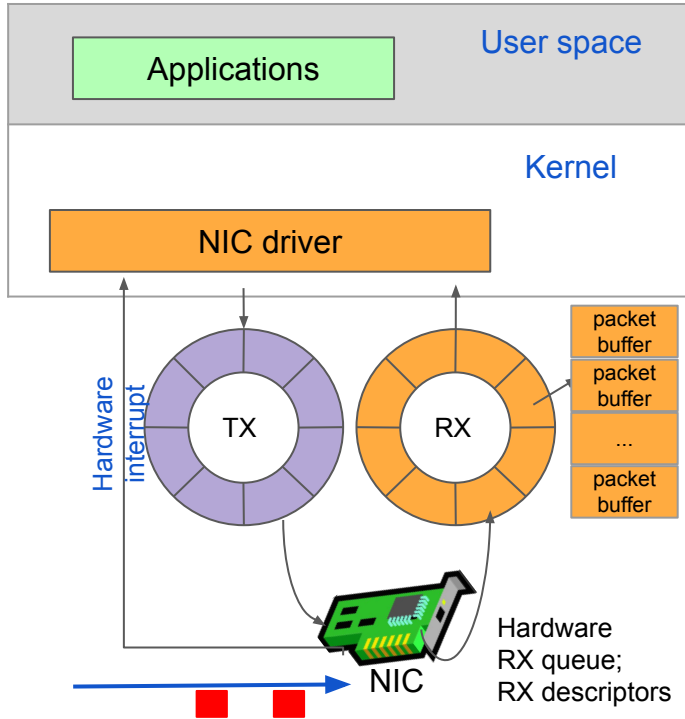


How does the application send and receive packets?



The journey of a packet through the Linux network stack ...

RX path: Packet arrives at the destination NIC



NIC receives the packet

- Match destination MAC address
- Verify Ethernet checksum (FCS)

Packets accepted at the NIC

- DMA the packet to RX ring buffer
- NIC triggers an interrupt


TX/RX rings

- Circular queue
- Shared between NIC and NIC driver
- Content: Length + packet buffer pointer

Interrupt processing in the linux kernel

- Top-half
 - Minimal processing
- Bottom-half
 - Rest of interrupt processing

Top-half interrupt processing



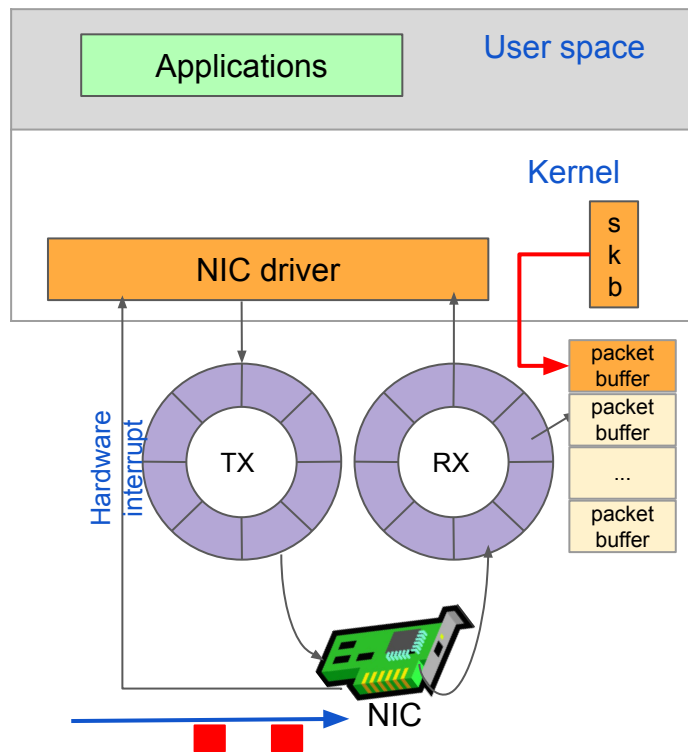
CPU interrupts the process in execution

Switch from user space to kernel space

Top-half interrupt processing

- Lookup IDT (Interrupt Descriptor Table)
- Call corresponding ISR (Interrupt Service Routine)
 - Acknowledge the interrupt
 - Schedule bottom-half processing
- Switch back to user space

Bottom-half processing



CPU initiates the bottom-half when it is free (soft-irq)

Switch from user space to kernel space

Driver dynamically allocates an **sk_buff** (a.k.a., skb)

Oops!!

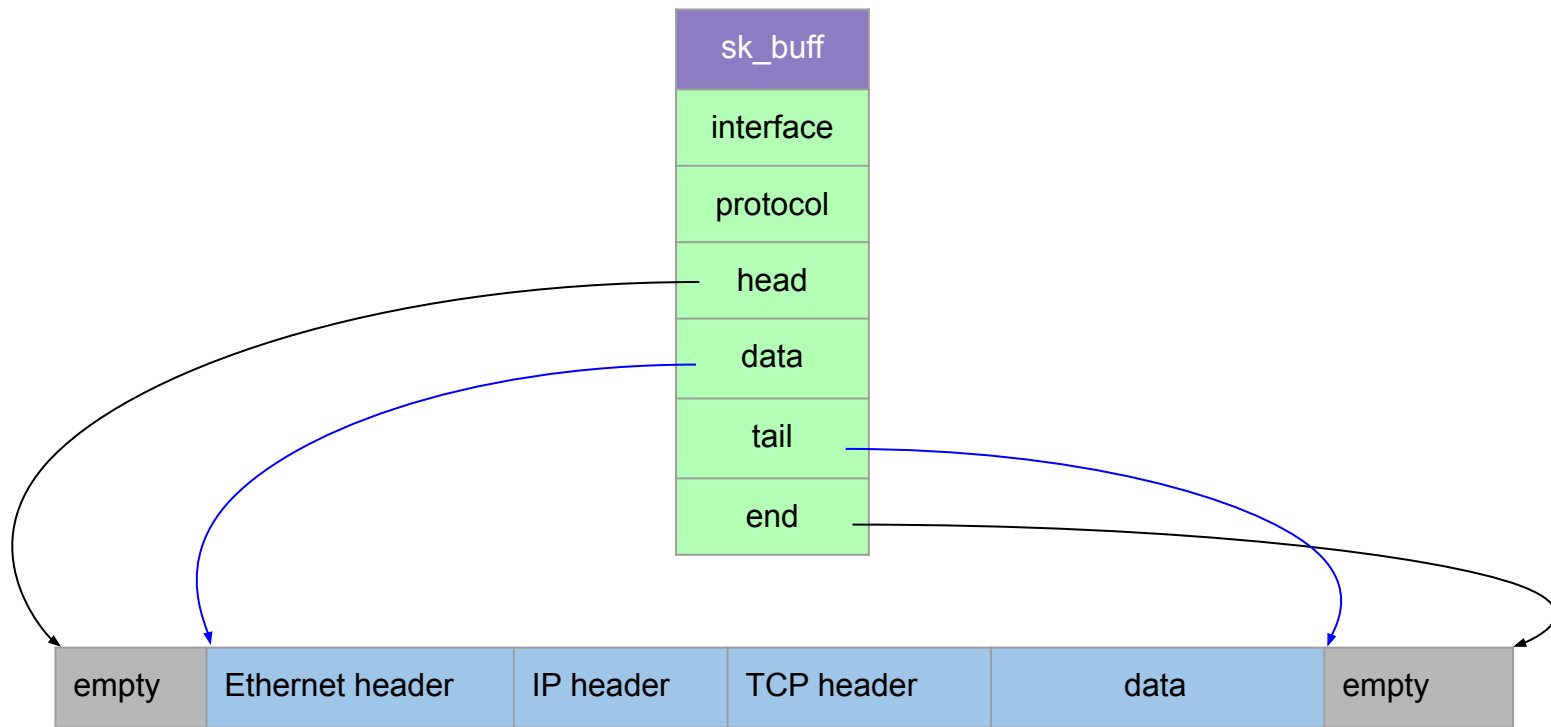
sk_buff ([sk_buff tutorial link](#))

In-memory data structure that contains packet metadata

- Pointers to packet headers and payload
- More packet related information ...

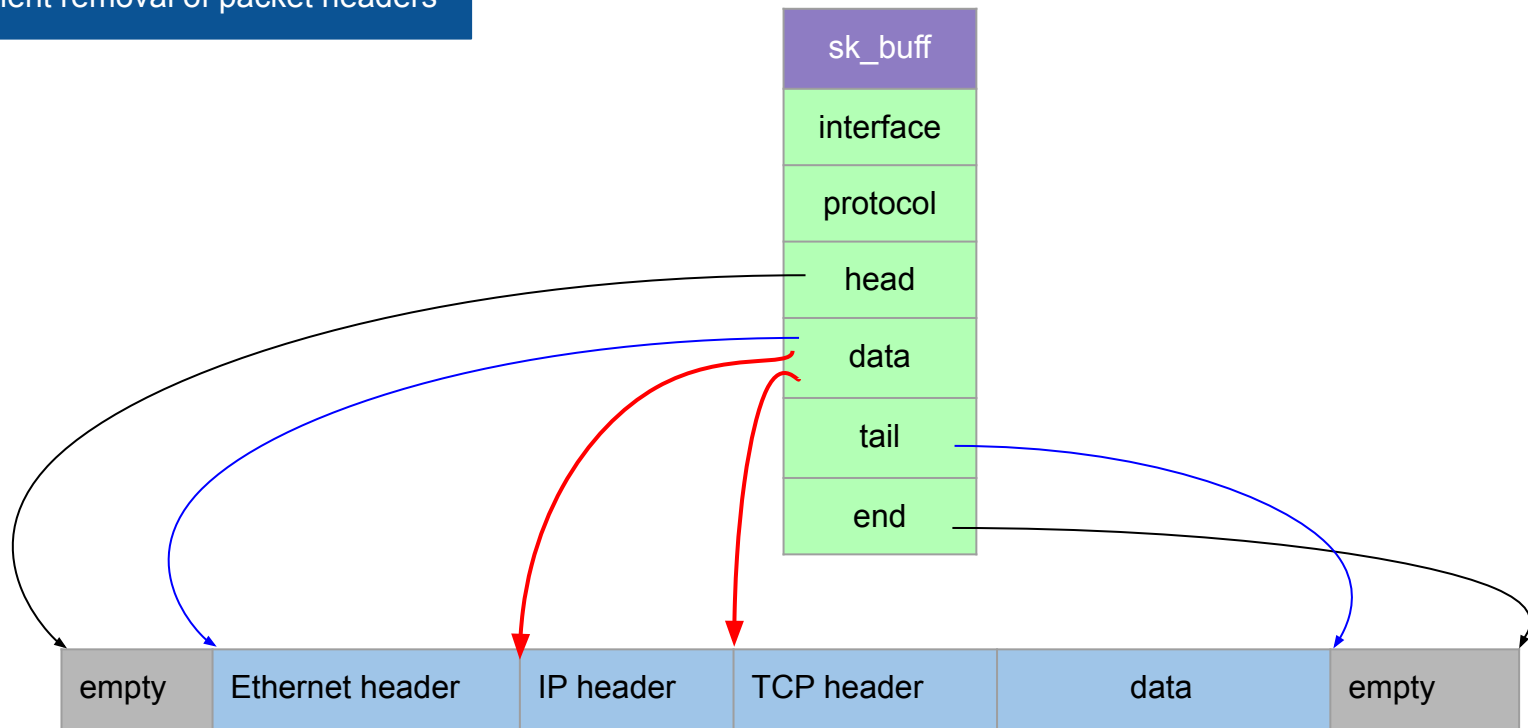
What is “sk_buff” ?

sk_buff: in-memory data structure that contains packet metadata



Advantages of sk_buff

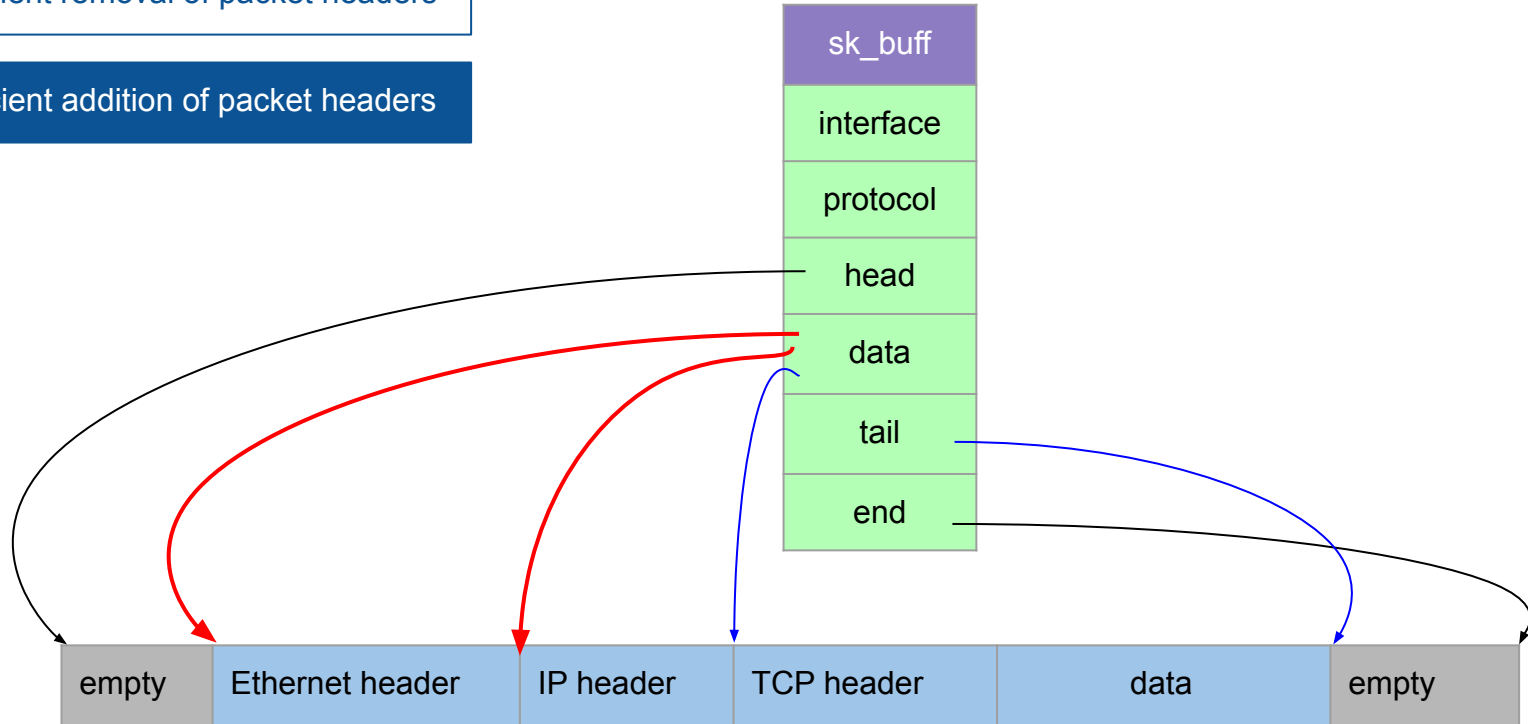
Efficient removal of packet headers



Advantages of sk_buff

Efficient removal of packet headers

Efficient addition of packet headers

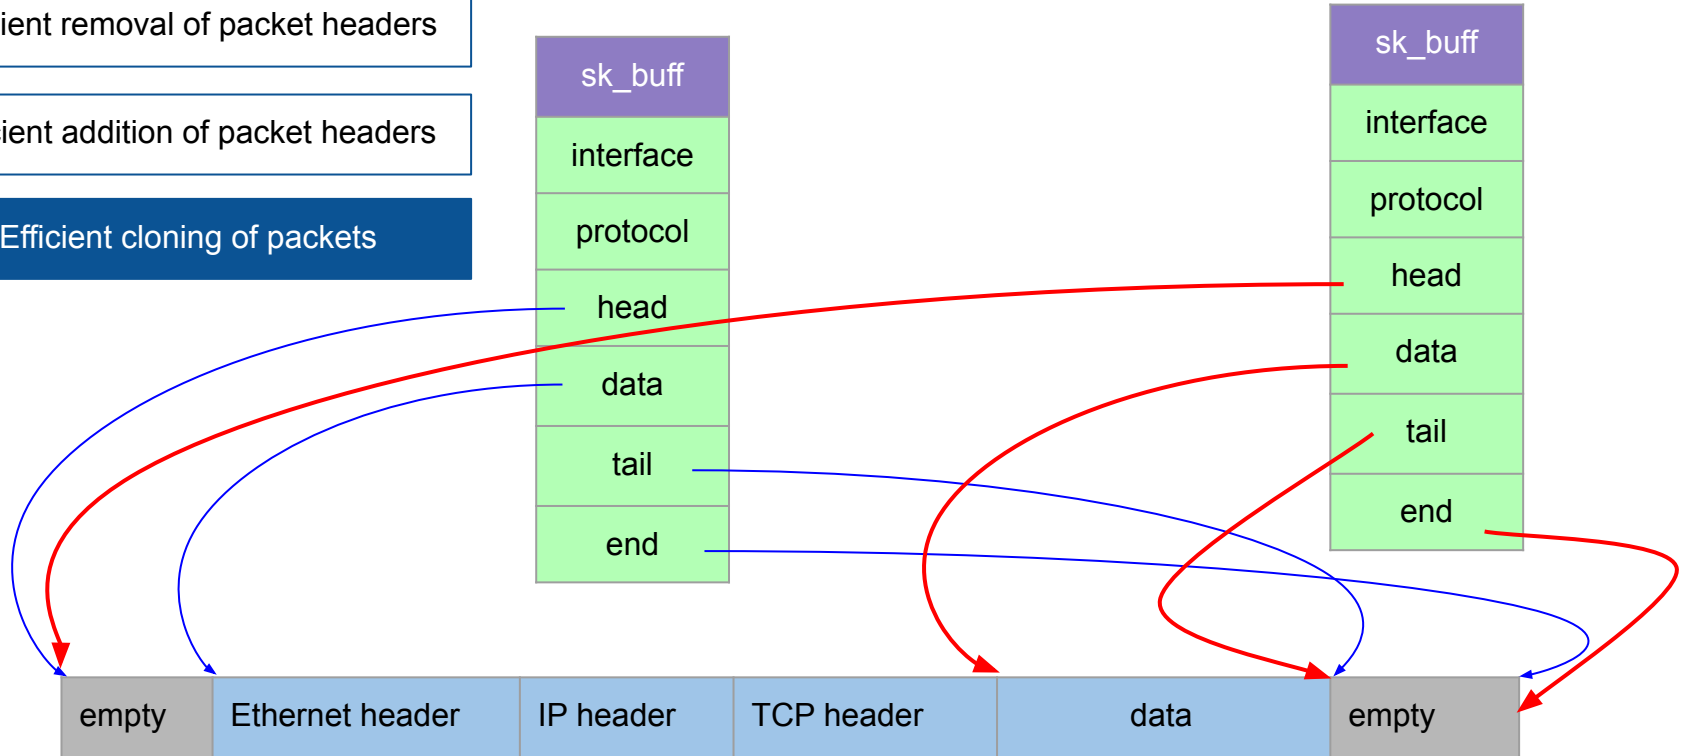


Advantages of sk_buff

Efficient removal of packet headers

Efficient addition of packet headers

Efficient cloning of packets

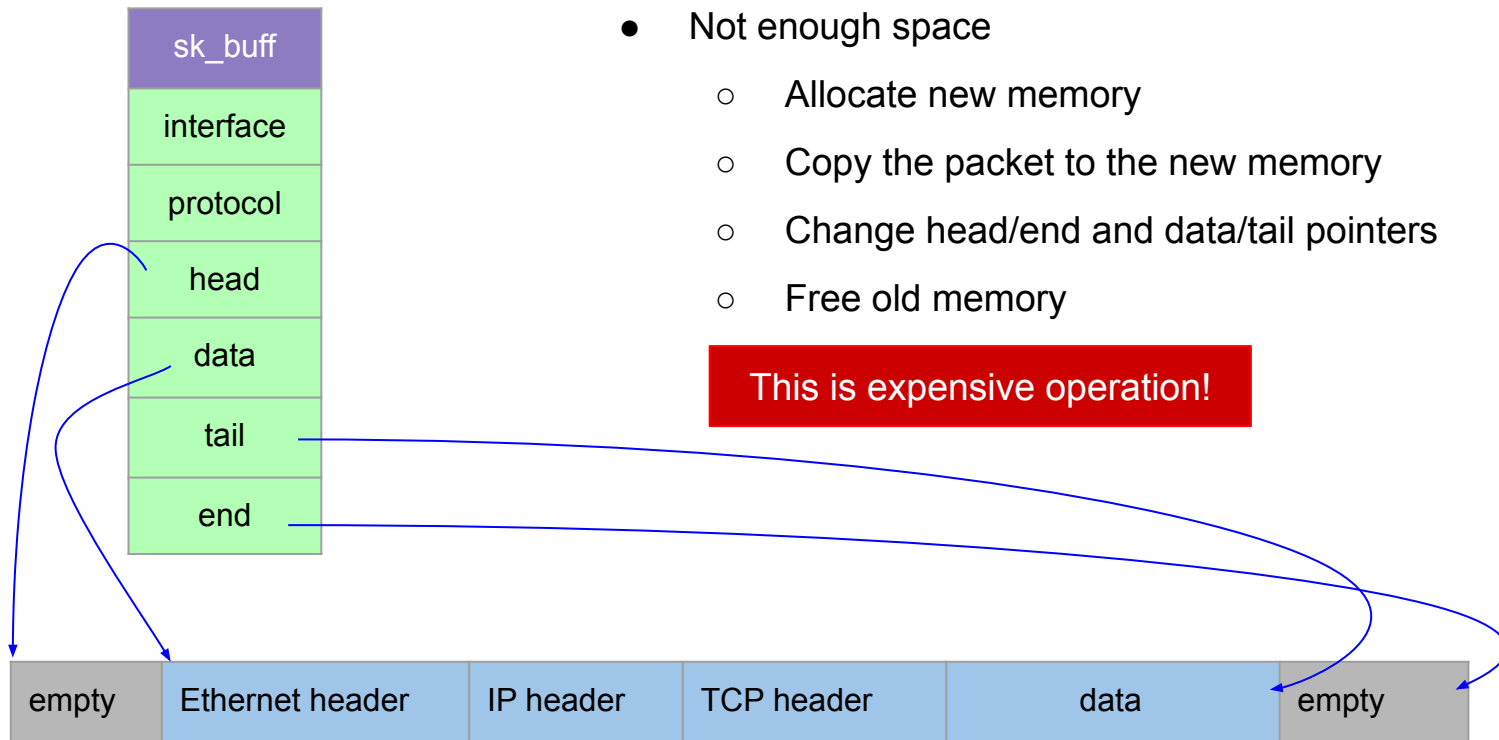


Disadvantage of sk_buff


Add packet header at the beginning --- encapsulation

- Not enough space
 - Allocate new memory
 - Copy the packet to the new memory
 - Change head/end and data/tail pointers
 - Free old memory

This is expensive operation!



More about skb structure



- More pointers
 - mac-header
 - net-header
 - transport-header
- Many many more fields ...

`sk_buff` never forgets!

After alloc_skb



After reserve_skb



skb containing data



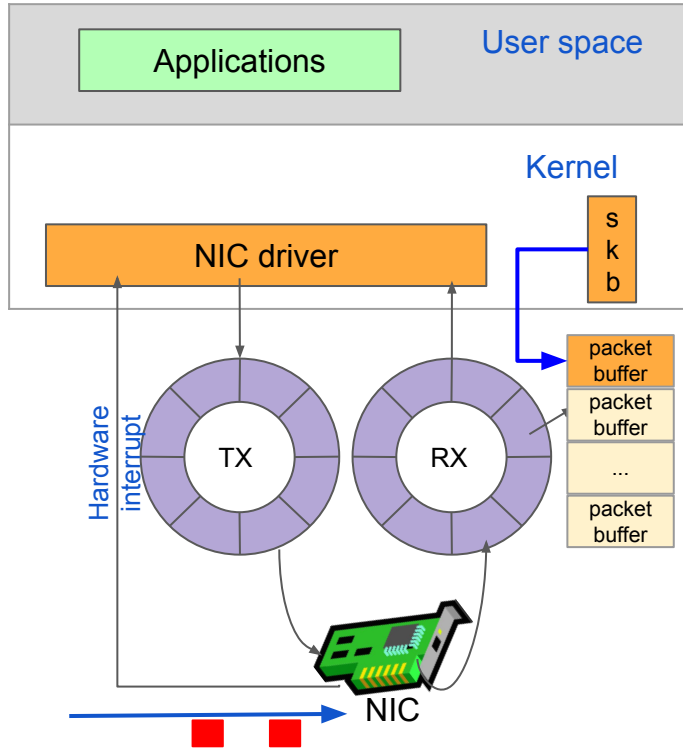
skb_put called on buffer



Skb_push has occurred on previous buffer



Bottom-half processing



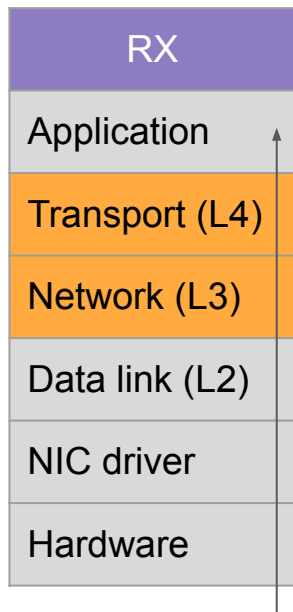
NIC driver processing

For all packets
in buffer

1. Driver dynamically allocates an **sk_buff**
2. Update sk_buff with packet metadata
3. Remove the Ethernet header
4. Pass sk_buff to the network stack

Call L3 protocol handler

L3/L4 processing



Common processing

1. Match destination IP/socket
2. Verify checksum
3. Remove header

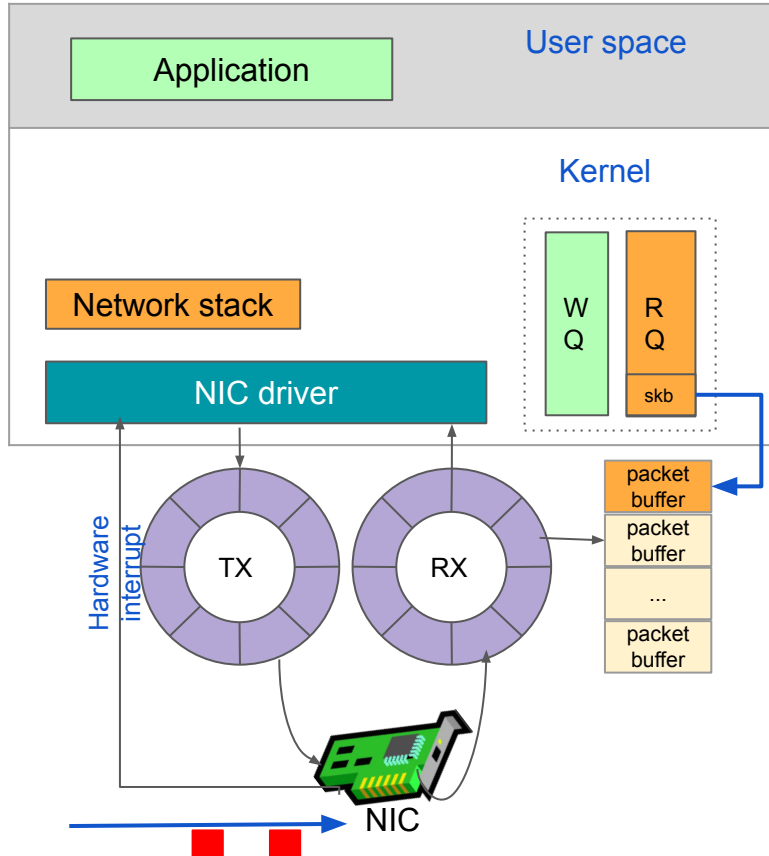


L3-specific processing

1. Route lookup
2. Combine fragmented packets
3. Call L4 protocol handler

L4-specific processing

L3/L4 processing



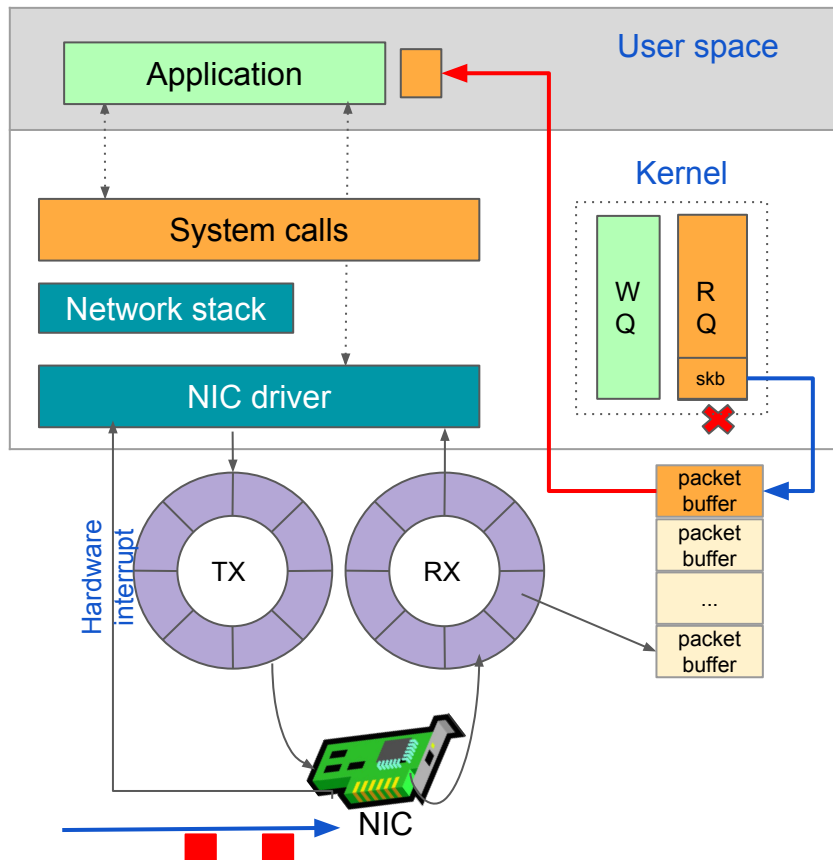
L3-specific processing

1. Route lookup
2. Combine fragmented packets
3. Call L4 protocol handler

L4-specific processing

1. Handle TCP state machine
2. Enqueue to socket read queue
3. Signal the socket

Application processing



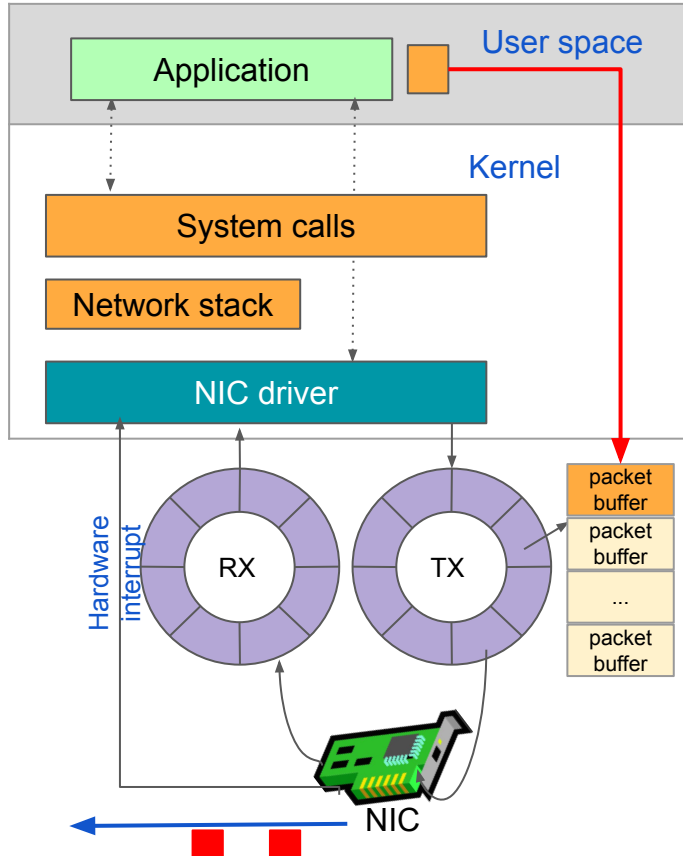
On socket read:

user space to kernel space

- Dequeue packet from socket receive queue (kernel space)
- Copy packet to application buffer (user space)
- Release sk_buff
- Return back to the application

kernel space to user space

Transmit path of an application packet

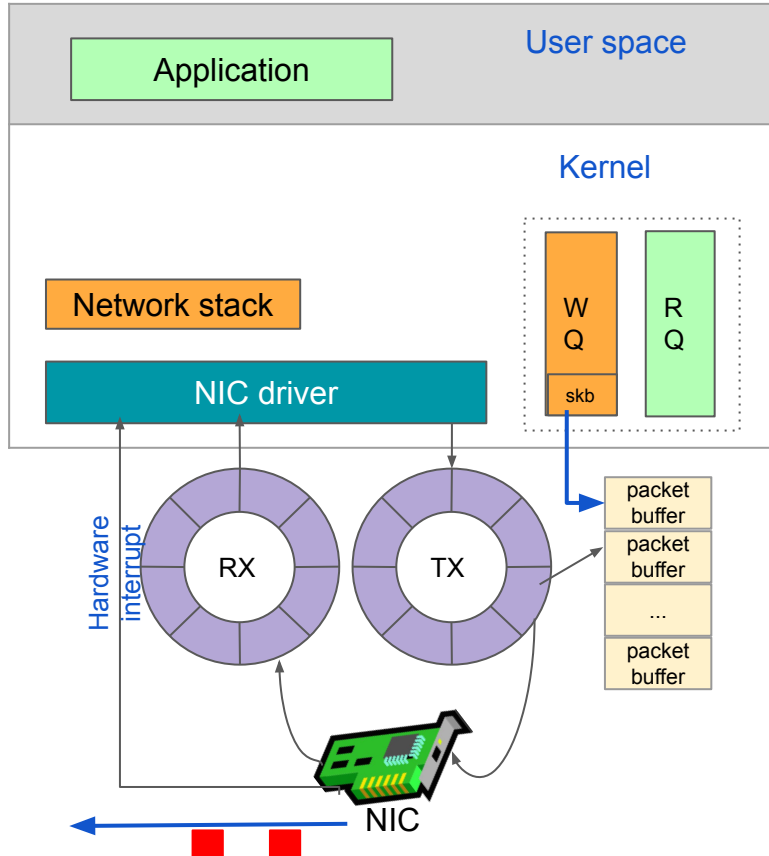


On socket write:

user space to kernel space

- Writes the packet to the kernel buffer
- Calls socket's send function (e.g., sendmsg)

L4/L3 processing



L4-specific processing

1. Allocate `sk_buff`
2. Enqueue `sk_buff` to socket write queue
3. Call L3 protocol handler

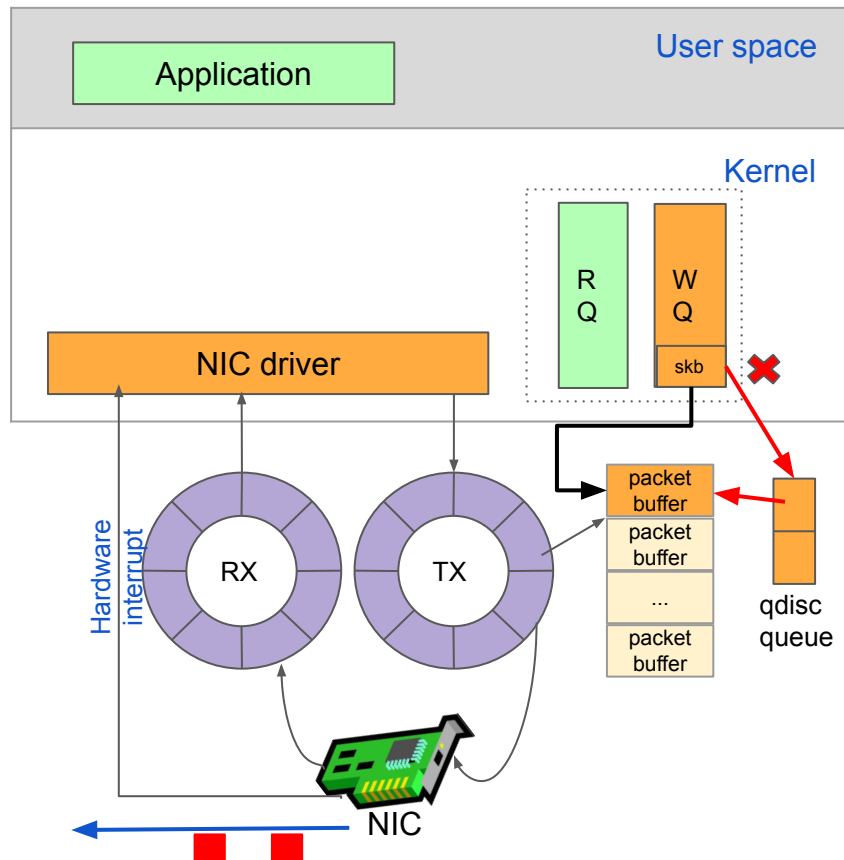
Common processing

1. Build header
2. Add header to packet buffer
3. Update `sk_buff`

L3-specific processing

1. Fragment, if needed
2. Call L2 protocol handler

L2 processing



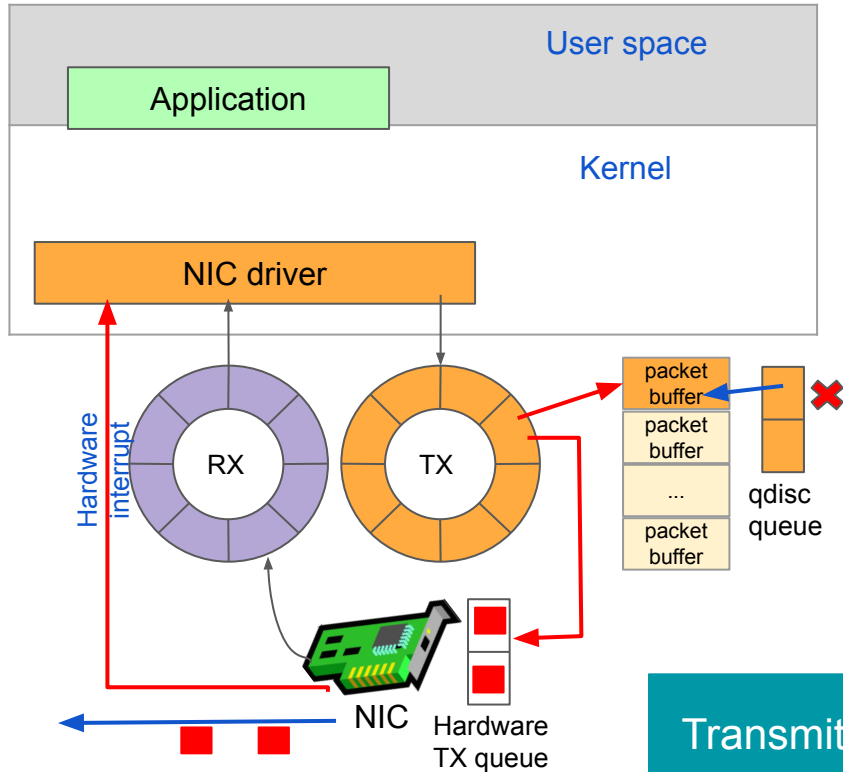
Enqueue packet to **queue discipline** (qdisc)

- Hold packets in a queue
- Apply scheduling policies (e.g. FIFO, priority)

qdisc

- Dequeue sk_buff (if NIC has free buffers)
- Post process sk_buff
 - Calculate IP/TCP checksum
 - ... (tasks that h/w cannot do)
- Call NIC driver's send function

NIC processing



NIC driver

- If hardware transmit queue full
 - Stop qdisc queue
- Otherwise:
 - Map packet data for DMA
 - Tells NIC to send the packet

NIC

- Calculates ethernet frame checksum (FCS)
- Sends packet to the wire
- **Sends an interrupt “Packet is sent” (kernel space to user space)**
- Driver frees the `sk_buff`; starts the qdisc queue

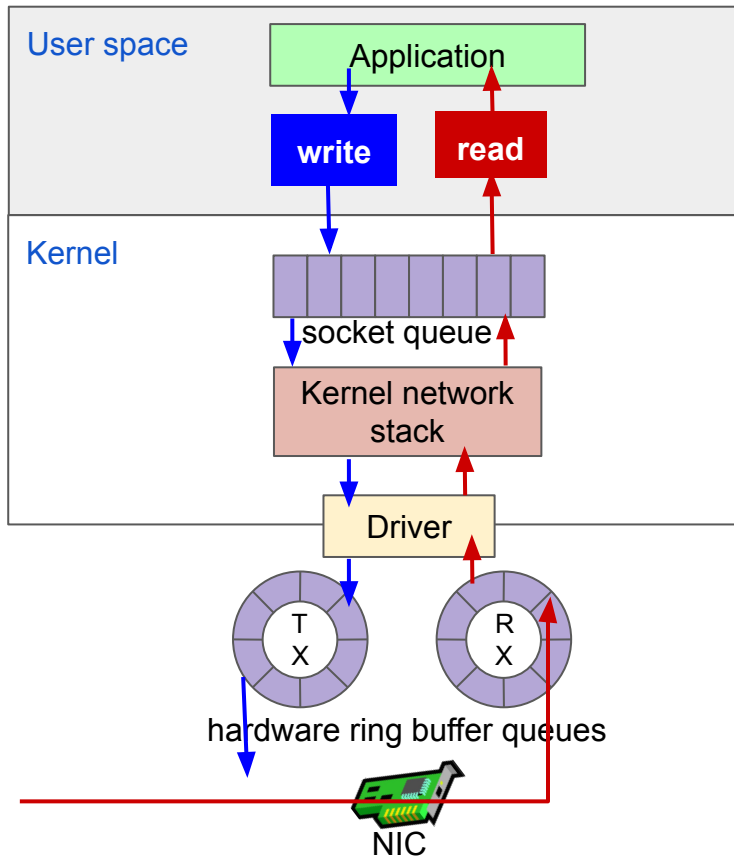
Transmit and receive packet processing pipeline DONE!!

Packet processing overheads in the kernel

- Too many context switches!!
 - Pollutes CPU cache
- Per-packet interrupt overhead
- Dynamic allocation of `sk_buff`
- Packet copy between kernel and user space
- Shared data structures

Cannot achieve line-rate for recent high speed NICs!! (40Gbps/100Gbps)

Packet processing overheads with the TCP/IP network stack



Sources of overhead

Mode switching

Context switching

Lock/unlock

Packet copy

sk_buff's dynamic
alloc/dealloc

Per packet interrupts

100 Gbps NIC¹

- RX: 20 CPUs
- TX: 10 CPUs



[1] [Understanding Host Network Stack Overheads](#)

* Slide abstracts details for simplicity

NIC Offloads? Why?

- Free up server CPU cycles for application
- Specialized processing can be efficient
- Scaling performance
 - Low latency
 - High throughput
- Power savings

Evolution of Network Interface Cards

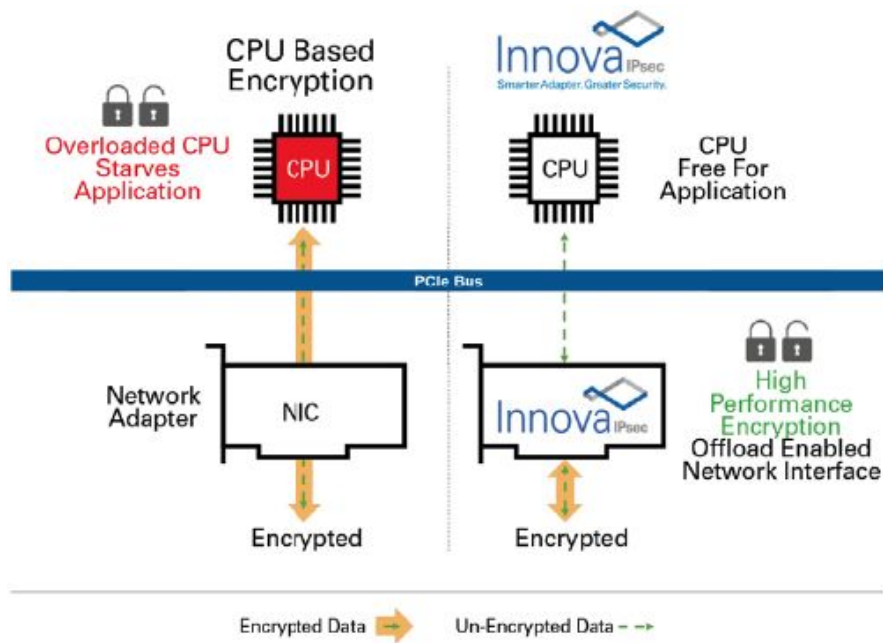
Improve NIC performance: What should be offloaded?

Basic offloads

- High frequency, specialized, compute-intensive tasks **Checksum offload**
- Reduce per packet overheads → Increase packet sizes?? **Segmentation offload**
- Parallel packet processing **Multiqueue offload**

Advanced offloads: Data plane in Hardware

Fixed function (minimally configurable) pipeline



- 6x throughput
- 10x CPU savings

Advanced offloads: Data plane in Hardware

Programmable pipeline

- Network processing unit (NPU)
 - Multithreaded or many-cores
 - Some domain-specific instructions
- FPGA
 - Gate-level programmable
- General purpose processor





Netronome smart NIC CX4000

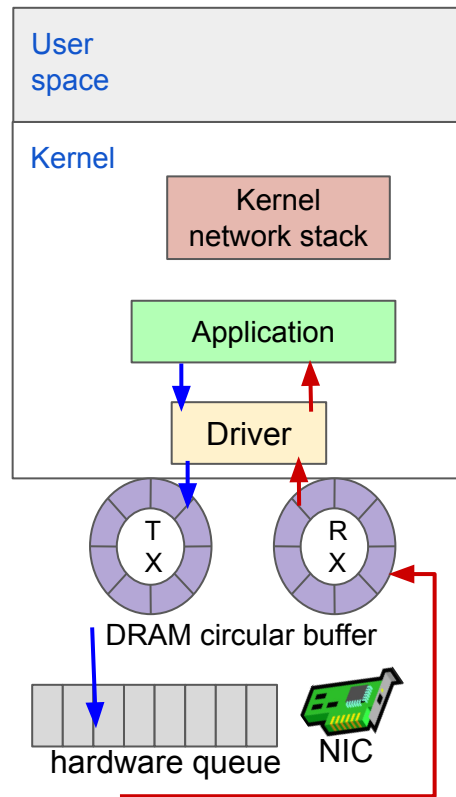


[Xilinx SN1022 FPGA smart NIC](#)

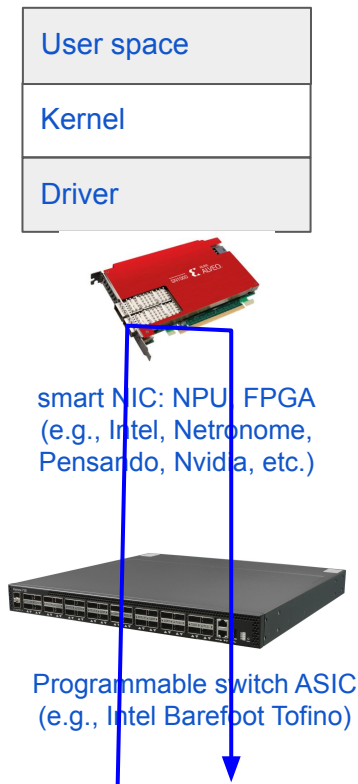
Evolution of network packet processing



Kernel bypass
(e.g. netmap, DPDK)




in-kernel compute
(e.g. eBPF, XDP)



in-network compute
(NIC/switch)

* Slide abstracts details for simplicity

Industry trends



Intel Mounts Evan IPU (ASIC+FPGA)



Pensando DPU DSC-200



Xilinx Alveo 280 FPGA smart NIC

Other vendors: Broadcom, Nvidia, ...

Programmable Networking Lab @ IIIT Delhi


<https://github.com/pnl-iiitd>



Intel Tofino switch



**Netronome smart NIC
CX4000**



**Xilinx MPSoC FPGA board
ZCU106**



**Xilinx FPGA smart NIC
SN1022**

Backup slides

After alloc_skb



After reserve_skb



skb containing data



skb_put called on buffer



Skb_push has occurred on previous buffer

